

Anly590_HW1_ArcherA

October 14, 2018

1 Anly 590 HW 1

1.1 Alex Archer

1.1.1 (1) Feedforward

1. Neural Network Diagram

$$\begin{aligned} 2. \quad z_i &= V_{i-1,0} \text{relu}(W_{0,0}^{(2)} \text{relu}(W_{0,0}^{(1)} x_1 + W_{0,1}^{(1)} x_2 + b_0^{(1)}) + W_{0,1}^{(2)} \text{relu}(W_{1,0}^{(1)} x_1 + W_{1,1}^{(1)} x_2 + b_1^{(1)}) + \\ &W_{0,2}^{(2)} \text{relu}(W_{2,0}^{(1)} x_1 + W_{2,1}^{(1)} x_2 + b_2^{(1)}) + b_0^{(2)}) + V_{i-1,1} \text{relu}(W_{1,0}^{(2)} \text{relu}(W_{0,0}^{(1)} x_1 + W_{0,1}^{(1)} x_2 + b_0^{(1)}) + \\ &W_{1,1}^{(2)} \text{relu}(W_{1,0}^{(1)} x_1 + W_{1,1}^{(1)} x_2 + b_1^{(1)}) + W_{1,2}^{(2)} \text{relu}(W_{2,0}^{(1)} x_1 + W_{2,1}^{(1)} x_2 + b_2^{(1)}) + b_1^{(2)}) + c_{i-1}, i = 1, 2, 3 \\ &\text{where } \text{relu}(x) = \max(0, x) \\ &\hat{y} = \text{softmax}(z) = \sigma(z), \text{ where } z = [z_1, z_2, z_3] \\ &\hat{y} = \sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^3 e^{z_k}} \text{ for } j = 1, 2, 3 \end{aligned}$$

3.


```
In [906]: import numpy as np
          import matplotlib.pyplot as plt
          import pandas as pd

In [907]: # activation functions

          def relu(x):
              return np.maximum(x, 0)

          def softmax(z):
              return np.exp(z)/(np.sum(np.exp(z),axis=0))

In [908]: # weights, inputs
          W1 = np.array([[1,0],[-1,0],[0,.5]])
          W2 = np.array([[1,0,0],[-1,-1,0]])
          V = np.array([[1,1],[0,0],[-1,-1]])
          b1 = np.array([0,0,1])
          b2 = np.array([1,-1])
          c = np.array([1,0,0])
          X = np.array([[1,0,0],[-1,-1,1]])
```



In [910]: # define the network

```
def ff_nn_2_ReLu(W1, W2, V, b1, b2, c, X):
    #z = softmax(np.dot(V,relu(np.dot(W2,relu(np.dot(W1,X)+b1)) + b2))+c)
    a1 = W1.dot(X) + b1.reshape(3,1)
    H1 = relu(a1)
    a2 = W2.dot(H1) + b2.reshape(2,1)
    H2 = relu(a2)
    a3 = V.dot(H2) + c.reshape(3,1)
    z = softmax(a3)
    return(z)
```

4.

In [911]: # weights

```
ff_nn_2_ReLu(W1,W2,V,b1,b2,c,X)
```


Out [911]: array([[0.94649912, 0.84379473, 0.84379473],
[0.04712342, 0.1141952 , 0.1141952],
[0.00637746, 0.04201007, 0.04201007]])

1.1.2 (2) Gradient Descent

1. $f(x,y) = (1-x)^2 + 100(y-x^2)^2$
 $\frac{\partial f(x,y)}{\partial x} = -2(1-x) - 400x(y-x^2)$
 $\frac{\partial f(x,y)}{\partial y} = 200(y-x^2)$

2.

```
In [912]: # contour plot
delta = .1
x = np.arange(-1.5, 1.5, delta)
y = np.arange(-1.0, 2, delta)
X, Y = np.meshgrid(x, y)
Z = (1-X)**2 + 100*(Y-X**2)**2
fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
```



3.

```
In [913]: # grad func
def grad_f(vector):
    x, y = vector
    df_dx = -2 * (1-x) - 400*x * (y-x**2)
    df_dy = 200 * (y-x**2)
    return np.array([df_dx, df_dy])

In [914]: # grad descent w/o momentum
def grad_descent(starting_point=None, iterations=100, learning_rate=.1):
    if starting_point:
        point = starting_point
    else:
        point = np.random.uniform(-1,1,size=2)
```

```

trajectory = [point]

for i in range(iterations):
    grad = grad_f(point)
    point = point - learning_rate * grad
    trajectory.append(point)
    #print(point)
return np.array(trajectory)

```

```

In [915]: np.random.seed(10)
          traj = grad_descent(iterations=100, learning_rate=.005)

```

```


fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y,'-o')

```

```

Out[915]: [<matplotlib.lines.Line2D at 0x124ae4908>]

```



```

In [916]: traj = grad_descent(iterations=100, learning_rate=.001)


```

```

fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y,'-o')

```


Out [916]: [<matplotlib.lines.Line2D at 0x124c090f0>]



```
In [917]: traj = grad_descent(iterations=100, learning_rate=.0001)
```

```
fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y, '-o')
```

Out [917]: [<matplotlib.lines.Line2D at 0x124d29e80>]



4.


```
In [918]: # grad descent with momentum
def grad_descent_with_momentum(starting_point=None, iterations=10, alpha=.9, epsilon=0.0005):
    if starting_point:
        point = starting_point
    else:
        point = np.random.uniform(-10,10,size=2)
    trajectory = [point]
    v = np.zeros(point.size)

    for i in range(iterations):
        grad = grad_f(point)
        v = alpha*v + epsilon*grad
        point = point - v
        #print(point)
        trajectory.append(point)
    return np.array(trajectory)

In [921]: traj = grad_descent_with_momentum(iterations=100, epsilon=.0005, alpha=.025)

fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y, '-o')
```


Out [921]: [<matplotlib.lines.Line2D at 0x124ebb4a8>]



```
In [926]: traj = grad_descent_with_momentum(iterations=100, epsilon=.00025, alpha=.025)
```

```
fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y, '-o')
```


Out [926]: [<matplotlib.lines.Line2D at 0x12531cf60>]



```
In [928]: traj = grad_descent_with_momentum(iterations=100, epsilon=.00005, alpha=.05)
```

```
fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y,'-o')
```

```
Out[928]: [<matplotlib.lines.Line2D at 0x1254a2f28>]
```

1.1.3 (3) Backprop

$$\begin{aligned}
 1. \quad \delta^L &= \frac{\partial L}{\partial a^L} \sigma'(z^L) = \vec{y} - \vec{y} \\
 \frac{\partial L}{\partial V} &= \delta^L a^{L-1} = (\vec{y} - \vec{y}) \cdot out_2^T \\
 \frac{\partial L}{\partial c} &= \frac{\partial L}{\partial a^L} = \Sigma(\vec{y} - \vec{y}) \\
 \frac{\partial L}{\partial W^{(2)}} &= \delta^{L-1} a^{L-2} = V^T \cdot (\vec{y} - \vec{y}) \odot \frac{\partial}{\partial in_2} \text{relu}(in_2) \cdot out_1^T \\
 \frac{\partial L}{\partial b^{(2)}} &= \delta^{L-1} = \Sigma(V^T \cdot (\vec{y} - \vec{y})) \odot \frac{\partial}{\partial in_2} \text{relu}(in_2) \\
 \frac{\partial L}{\partial W^{(1)}} &= \delta^{L-2} X = \Sigma W^{(2)T} \cdot (V^T \cdot (\vec{y} - \vec{y})) \odot \frac{\partial}{\partial in_2} \text{relu}(in_2) \odot \frac{\partial}{\partial X} \text{relu}(X) \cdot X^T \\
 \frac{\partial L}{\partial b^{(1)}} &= \delta^{L-2} X = \Sigma W^{(2)T} \cdot (V^T \cdot (\vec{y} - \vec{y})) \odot \frac{\partial}{\partial in_2} \text{relu}(in_2) \odot \frac{\partial}{\partial X} \text{relu}(X) \\
 \text{where } \text{relu}(x) &= \max(x, 0), in_1 = W^{(1)} \cdot x + b^{(1)}, in_2 = W^{(2)} \cdot out_1 + b^{(2)}, out_1 = \text{relu}(in_1), out_2 = \text{relu}(in_2) \\
 \frac{\partial}{\partial x} \text{relu}(x) &= \begin{cases} 1, & x > 0 \\ 0, & \text{otherwise} \end{cases}
 \end{aligned}$$

In [1941]: # new grad func for the NN from part 1

```

def grad_f(params, x, y):
    # params
    W1 = params[0:6].reshape(3,2)
    W2 = params[6:12].reshape(2,3)
    V = params[12:18].reshape(3,2)
    b1 = params[18:21].reshape(3,1)
    b2 = params[21:23].reshape(2,1)
    c = params[23:26].reshape(3,1)

```

```

# forward pass
a1 = W1.dot(x) + b1
H1 = relu(a1)
a2 = W2.dot(H1) + b2
H2 = relu(a2)
a3 = V.dot(H2) + c
y_hat = softmax(a3)

# gradients
d_V = (y_hat - y).dot(H2.T)
d_c = (y_hat - y).sum(axis=1)
d_W2 = ((V.T.dot((y_hat - y))) * (H2 > 0)).dot(H1.T)
d_b2 = ((V.T.dot((y_hat - y))) * (H2 > 0)).sum(axis=1)
d_W1 = (W2.T.dot((V.T.dot((y_hat - y))) * (H2 > 0)) * (H1 > 0)).dot(x.T)
d_b1 = (W2.T.dot((V.T.dot((y_hat - y))) * (H2 > 0)) * (H1 > 0)).sum(axis=1)

# return vector of gradients
grad_vec = np.array([d_V, d_c, d_W2, d_b2, d_W1, d_b1])
grad_vec = [p.flatten() for p in grad_vec]
grad_vec = np.concatenate(grad_vec)
return(grad_vec)

```

```

In [1913]: def gen_gmm_data(n = 999, plot=False):
# Fixing seed for repeatability
np.random.seed(123)

# Parameters of a normal distribuion
mean_1 = [0, 2] ; mean_2 = [2, -2] ; mean_3 = [-2, -2]
mean = [mean_1, mean_2, mean_3] ; cov = [[1, 0], [0, 1]]

# Setting up the class probabilities
n_samples = n
pr_class_1 = pr_class_2 = pr_class_3 = 1/3.0
n_class = (n_samples * np.array([pr_class_1, pr_class_2, pr_class_3])).astype(int)

# Generate sample data
for i in range(3):
    x1,x2 = np.random.multivariate_normal(mean[i], cov, n_class[i]).T
    if (i==0):
        xs = np.array([x1,x2])
        cl = np.array([n_class[i]*[i]])
    else:
        xs_new = np.array([x1,x2])
        cl_new = np.array([n_class[i]*[i]])
        xs = np.concatenate((xs, xs_new), axis = 1)
        cl = np.concatenate((cl, cl_new), axis = 1)

```

```

        # One hot encoding classes
        y = pd.Series(cl[0].tolist())
        y = pd.get_dummies(y).as_matrix()

        # Normalizing data (prevents overflow errors)
        mu = xs.mean(axis = 1)
        std = xs.std(axis = 1)
        xs = (xs.T - mu) / std

    return xs, y, cl

```


```

In [2225]: x,y,cl = gen_gmm_data(500)
           plt.scatter(x[:,0], x[:,1], c=cl.reshape(498,))

```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:28: FutureWarning: Method .as_matrix() is deprecated, use .to_matrix() instead.

Out[2225]: <matplotlib.collections.PathCollection at 0x14039e550>



```

In [1915]: def loss(y, y_hat):
            # cross entropy
            tot = y * np.log(y_hat)
            return(-tot.sum())

            def yhat(in_vec, params):
                # yhat helper

```

```

w_1 = params[0:6].reshape(3,2)
w_2 = params[6:12].reshape(2,3)
v = params[12:18].reshape(3,2)
b_1 = params[18:21].reshape(3,1)
b_2 = params[21:23].reshape(2,1)
c_0 = params[23:26].reshape(3,1)

a_1 = w_1.dot(in_vec) + b_1
H_1 = relu(a_1)
a_2 = w_2.dot(H_1) + b_2
H_2 = relu(a_2)
a_3 = v.dot(H_2) + c_0
y_hat = softmax(a_3)

return(y_hat)

```

```

In [2296]: # grad descent w/o momentum for NN in part 1
def grad_descent(x, y, iterations=10, learning_rate=.01):
    point = np.random.uniform(-.1, .1, size = 26).astype("float128")

    trajectory = [point]
    losses = [loss(y, yhat(x, point))]

    for i in range(iterations):
        grad = grad_f(point, x, y)
        point = point - learning_rate * grad
        trajectory.append(point)
        losses.append(loss(y, y_hat(x, point)))
    return(np.array(trajectory), losses)

```

```

In [2980]: traj, losses = grad_descent(x.T, y.T, iterations=350,
                                         learning_rate=.00025)

```

```


In [2981]: plt.plot(losses)

```

```

Out[2981]: [<matplotlib.lines.Line2D at 0x1512875c0>]

```




```
In [2773]: # gradient descent with momentum for NN in part 1
def grad_descent_with_momentum2(x, y, iterations=100,
                                alpha=.75, epsilon=.1):
    point = np.random.uniform(-.1, .1, size = 26).astype("float128")
    v = np.zeros(point.size)
    trajectory = [point]
    losses = [loss(y, yhat(x, point))]

    for i in range(iterations):
        grad = grad_f(point, x, y)
        v = alpha*v + epsilon*grad
        point = point - v
        trajectory.append(point)
        losses.append(loss(y, y_hat(x, point)))
    return(np.array(trajectory), losses)

In [2769]: traj_m, losses_m = grad_descent_with_momentum2(x.T, y.T, iterations=350,
                                                           alpha=.001, epsilon=.0005)

In [2770]: plt.plot(losses_m)

Out[2770]: [<matplotlib.lines.Line2D at 0x14c3ea128>]
```



The gradient descent with momentum loss appears to be similar to gradient descent without momentum. They converge at a similar rate. The parameters likely need to be tuned a bit more to find the optimal results.